AI Ethics

A Reading List

Curated for the pivot into AI ethics & governance — grouped by purpose, not alphabetically.

How to read this list. These are grouped into four shelves. You don't need all of them, and you shouldn't start at the top and grind down. A suggested order is at the end. The honest note under each tells you what it's actually good for — and, where relevant, its blind spot — so you can pick by what you need rather than by hype.

1 · Foundations

The core map — read one or two of these first; they give you the vocabulary everything else assumes.

The Alignment Problem — *Brian Christian*

The clearest single narrative of how AI systems misread their goals, act in unintended ways, and reinforce societal bias — told through real cases across healthcare, criminal justice, and hiring.

Why it's on your list: The best on-ramp. It connects the technical and the ethical without requiring you to be an engineer, and its bias/fairness chapters map directly onto the cases you've been analysing.

AI Ethics — *Mark Coeckelbergh (MIT Press Essential Knowledge)*

A short, precise survey of the whole field: responsibility, transparency, bias, autonomy, labour, privacy, and the limits of machine decision-making.

Why it's on your list: The cleanest starting point if you want the map before the deep dives — exactly the operational vocabulary the AIGP and your portfolio spine draw on. Small enough to read in a weekend.

Atlas of AI — *Kate Crawford*

Drags AI 'back down to earth' — the mining, energy, hidden human labour, and power structures behind systems we treat as abstract. A materialist counter-narrative to the 'AI as neutral cloud' story.

Why it's on your list: Sharpens the instinct you keep returning to — that harm sits in human choices and structures, not in the technology itself. The 'who bears the cost?' lens you used on Helsinki, argued at book length.

2 · Fairness, Bias & Discrimination

The area you've done the most work in — these deepen the exact themes from your peer reviews.

Weapons of Math Destruction — *Cathy O'Neil*

The classic on how opaque scoring algorithms scale injustice in lending, hiring, policing, and education — how a biased model becomes a self-reinforcing feedback loop.

Why it's on your list: This is your Helsinki 'high-risk drifts toward high-cost' and the Workday/COMPAS cases, written before they were headlines. Foundational for the fairness deck.

Unmasking AI — *Joy Buolamwini*

By the researcher whose 'Gender Shades' work exposed that facial recognition fails far more often on darker-skinned and female faces. Part memoir, part investigation.

Why it's on your list: The direct backstory to the Robert Williams wrongful-arrest case you just reviewed — the evidence that facial-recognition error falls along racial lines, which is what makes it discrimination and not just error.

Race After Technology — *Ruha Benjamin*

Introduces the 'New Jim Code' — how systems that appear neutral or benevolent can deepen discrimination, precisely because their neutrality is assumed.

Why it's on your list: The deep version of the proxy problem and 'formal fairness masking substantive discrimination' — the two lifts you kept flagging in others' bias answers.

3 · Governance, Law & Policy

Your actual destination. Lighter on philosophy, heavier on 'what do we implement?' — and doubles as AIGP prep.

The Oxford Handbook of Ethics of AI — *Dubber, Pasquale & Das (eds.)*

A comprehensive reference collection spanning ethics, law, and policy — the kind of volume you consult by chapter rather than read cover to cover.

Why it's on your list: A standing reference for the governance shelf. Use it to go deep on a specific question (accountability, transparency law, human rights) when a deck or essay needs it.

Robot Rules: Regulating Artificial Intelligence — *Jacob Turner*

One of the earliest serious treatments of how law can actually grip AI — liability, legal personality, and translating vague ethical principles into workable government mechanisms.

Why it's on your list: Directly on your enforceability instinct — the move from 'principles' to 'mechanisms that actually bind.' The sunset-clause and guardrail-enforceability questions, in legal form.

An EU AI Act practitioner's guide (choose a current edition) — *various — e.g. Kauffman / Soares-style controls guides*

Field-guide treatments of the EU AI Act for people who ship real products: risk classification, conformity assessments, human-oversight mechanisms, and how the AI Act and NIST AI RMF map onto each other.

Why it's on your list: The single most AIGP-relevant reading. Domain II of your syllabus is largely the EU AI Act + NIST + ISO — a good practitioner guide is exam prep and portfolio fuel at once. Pick the most recent edition, as these date quickly.

4 · Big-Picture & Perspective

Optional, but they widen the frame — useful for the 'why this matters' voice in your portfolio.

AI Snake Oil — *Arvind Narayanan & Sayash Kapoor*

A razor-sharp demystification of the AI hype cycle — separating what these systems genuinely do from what's oversold, especially predictive AI.

Why it's on your list: Trains the exact scepticism you've been practising — spotting when a benign-sounding claim ('just predict risk') hides a weaker or more dangerous reality. Excellent bias-detector for your own writing.

The Age of Surveillance Capitalism — *Shoshana Zuboff*

The landmark (and demanding) account of how behavioural data became the raw material of a new economic order — the deep background to every privacy and 'greater good' tracking question.

Why it's on your list: The book-length version of your pandemic-tracking 'greater good' deck: why data infrastructure, once built, outlives the emergency that justified it. Long — read selectively.

A suggested order

Not top-to-bottom — this sequence matches where you already are.

You've already done a lot of the thinking these books contain, so treat them as deepening and evidencing what you know — not as starting from zero.

Start: Coeckelbergh, AI Ethics — get the whole map cheaply and quickly, so the rest slots into a structure.

Then: Christian, The Alignment Problem — the narrative that connects the technical and ethical; the richest single read.

Your strong suit: O'Neil + Buolamwini — these evidence the fairness/bias work you're already best at, and back your deck and peer-review instincts with cases and data.

Toward the goal: A current EU AI Act practitioner guide + Turner, Robot Rules — the pivot to governance, and double-duty AIGP prep.

For voice & scepticism: Narayanan & Kapoor, AI Snake Oil, and dip into Crawford / Zuboff — perspective and the 'why it matters' register for your portfolio.

A note on freshness. The Foundations and Fairness shelves are durable — those ideas won't date. The Governance shelf moves fast: laws, editions, and frameworks change yearly, so always check for the newest edition of any practitioner guide, and treat the primary sources (the EU AI Act text, the NIST AI RMF, the OECD principles, ISO 42001) as the real ground truth that the books interpret.